

Rails Performance Best Practices

<http://ihower.tw>
2010/3

About Me

- 張文鈺 a.k.a. ihowr
 - <http://ihowr.tw>
 - <http://twitter.com/ihowr>
 - <http://github.com/ihowr>
- Ruby on Rails Developer since 2006
- Ruby Taiwan Community
 - <http://ruby.tw>



Performance Worse Practices

- Premature Optimization
- Guessing
- Caching everything
- Fighting the framework

Performance Guidelines

- Algorithmic improvements always beat code tweaks
- As a general rule, maintainability beats performance
- Only optimize what matters (80/20 rule)
- Measure twice, cut once
- Must be balanced with flexibility

How to improve performance?

- Find the target baseline
- Know where you are now
- Profile to find bottlenecks
- Remove bottlenecks
- Repeat

Agenda

- Analysis and Measurement
- Write Efficient Ruby code
- Use REE 1.8.7 or Ruby 1.9
- Use faster Ruby Library
- Caching
- SQL and ActiveRecord
- Consider NoSQL storage
- Rack and Rails Metal
- Use HTTP server for static files
- Front-end web performance
- Use external programs or write inline C code

Finding the Problems: Analysis

Server Log Analysis

- <http://github.com/wvanbergen/request-log-analyzer>
 - Request distribution per hour
 - Most requested
 - HTTP methods
 - HTTP statuses returned
 - Rails action cache hits
 - Request duration
 - View rendering time
 - Database time
 - Process blockers
 - Failed requests

Commercial Monitor Products

- New Relic
- Scout

Rack::Bug

Rails middleware which gives you an informative toolbar in your browser

The screenshot shows a web browser window titled "ASCIIcasts - 'Episode 159 - More on Cucumber'". The address bar shows the URL "http://localhost:3000/episodes/159-more-on-cucumber". The Rack::Bug toolbar is visible at the top of the page, displaying various metrics: Rack::Bug, Rails 2.3.2, 83.05ms, Request Vars, Rack Env, 4 Queries (7.00ms), 20 AR Objects, and Cache: 0.00ms (0 calls). Below the toolbar, the ASCIIcasts logo is displayed, followed by a navigation bar with links: Home, Newest episode, All episodes, and About ASCIIcasts. A search bar is also present. The main content area shows the title "159: More on Cucumber" with a link to the original Railscast. The text describes the episode's focus on Cucumber, a high-level testing framework, and mentions the use of Webrat and Factory Girl. A section titled "Updating Cucumber" discusses the recent update to version 0.3.0 and the need to update the application's configuration. A sidebar on the right lists various tags: active-record, active-support, administration, authentication, caching, controllers, forms, performance, plugins, rails-23, refactoring, security, and stuff.

ASCIIcasts - "Episode 159 - More on Cucumber"

http://localhost:3000/episodes/159-more-on-cucumber

RSS Google

links Toggle Rack::Bug

Rack::Bug Rails 2.3.2 83.05ms Request Vars Rack Env 4 Queries (7.00ms) 20 AR Objects Cache: 0.00ms (0 calls)

Templates: 18.29ms Log 176 KB Δ, 34644 KB total

> ASCIIcasts
video.to_s

Home Newest episode All episodes About ASCIIcasts

Search: search

159: More on Cucumber ([view original Railscast](#))

Episode 155 was all about [Cucumber](#), a high-level testing framework. This episode will show some of its more advanced features, so if you haven't seen the other episode then it's well worth [catching up](#) before reading this. We'll also be using Webrat and Factory Girl which were covered in episodes [156](#) and [158](#).

Updating Cucumber

Cucumber has recently been updated to version 0.3.0, so the first thing we'll need to do is update the version of the gem that our application's using. In `/config/environments/test.rb` we'll change the reference to Cucumber so that it uses the latest version.

Tags

- [active-record](#)
- [active-support](#)
- [administration](#)
- [authentication](#)
- [caching](#)
- [controllers](#)
- [forms](#)
- [performance](#)
- [plugins](#)
- [rails-23](#)
- [refactoring](#)
- [security](#)
- [stuff](#)

ASCIICasts - "Episode 159 - More on Cucumber"

http://localhost:3000/episodes/159-more-on-cucumber

RSS

Google

links Toggle Rack::Bug

Rack::Bug

Rails 2.3.2

83.05ms

Request Vars

Rack Env

4 Queries (7.00ms)

20 AR Objects

Cache: 0.00ms (0 calls)

Templates: 18.29ms

Log 176 KB Δ, 34644 KB total

Close

SQL Queries

Time (ms)	Query		
0.98ms	SELECT * FROM "episodes" WHERE ("episodes"."identifier" = '159-more-on-cucumber') LIMIT 1	Show Backtrace	SELECT EXPLAIN Profile
2.51ms	SELECT * FROM "episodes" WHERE (episode_id < 159) ORDER BY episode_id DESC LIMIT 1	Show Backtrace	SELECT EXPLAIN Profile
1.04ms	SELECT * FROM "episodes" WHERE (episode_id > 159) ORDER BY episode_id ASC LIMIT 1	Show Backtrace	SELECT EXPLAIN Profile
2.47ms	SELECT * FROM "tags" ORDER BY title ASC	Show Backtrace	SELECT EXPLAIN Profile

Open # on this page in a new tab

MemoryLogic

Adds in process id and memory usage in your rails logs, great for tracking down memory leaks

- <http://github.com/binarylogic/memorylogic>

```
Processing WelcomeController#index (for 127.0.0.1 at 2010-02-26 01:51:54) [GET] (mem 104792)
  Parameters: {"action"=>"index", "controller"=>"welcome"} (mem 104792)
```

```
...
```

```
Memory usage: 108888 | PID: 6170 (mem 108888)
```

```
Completed in 3570ms (View: 1861, DB: 1659) | 200 OK [http://localhost/] (mem 108888)
```

oink

Log parser to identify actions which significantly increase VM heap size

- <http://github.com/noahdl/oink>

```
script/oink -t 0 ~/rails-app/log/development.log
```

```
---- MEMORY THRESHOLD ----  
THRESHOLD: 0 MB
```

```
-- SUMMARY --
```

```
Worst Requests:
```

```
1. Feb 26 02:12:45, 37360 KB, SessionsController#new  
2. Feb 26 02:12:41, 37352 KB, BooksController#hot  
3. Feb 26 02:12:21, 16824 KB, BooksController#hot  
4. Feb 26 02:12:25, 11632 KB, BooksController#hot  
5. Feb 26 02:12:19, 11120 KB, BooksController#hot  
6. Feb 26 02:12:51, 9888 KB, WelcomeController#index  
7. Feb 26 02:12:28, 7548 KB, WelcomeController#index  
8. Feb 26 02:12:23, 5120 KB, ArticlesController#index
```

```
Worst Actions:
```

```
4, BooksController#hot  
2, WelcomeController#index  
1, ArticlesController#index  
1, SessionsController#new
```

ruby-prof gem

- a fast code profiler for Ruby. Its features include:
 - Speed - it is a C extension
 - Modes - call times, memory usage and object allocations.
 - Reports - can generate text and cross-referenced html reports
 - Threads - supports profiling multiple threads simultaneously
 - Recursive calls - supports profiling recursive method calls

ruby-prof example (I)

```
require 'ruby-prof'

# Profile the code
RubyProf.start
...
# code to profile
100.times { puts "blah" }

result = RubyProf.stop

# Print a flat profile to text
printer = RubyProf::FlatPrinter.new(result)
printer.print(STDOUT, 0)
```

ruby-prof example (2)

Thread ID: 2148368700

Total: 0.022092

%self	total	self	wait	child	calls	name
13.19	0.00	0.00	0.00	0.00	2	Readline#readline (ruby_runtime:0}
12.16	0.01	0.00	0.00	0.00	79	RubyLex#getc (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:101}
4.87	0.00	0.00	0.00	0.00	30	RubyLex#ungetc (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:144}
4.52	0.00	0.00	0.00	0.00	5	RubyLex#identify_identifier (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:770}
4.41	0.00	0.00	0.00	0.00	202	IO#write (ruby_runtime:0}
4.03	0.02	0.00	0.00	0.02	20	IRB::SLex::Node#match_io (/opt/ruby-enterprise/lib/ruby/1.8/irb/slex.rb:204}
4.00	0.01	0.00	0.00	0.01	28	Proc#call (ruby_runtime:0}
2.92	0.00	0.00	0.00	0.00	20	RubyToken#Token (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-token.rb:84}
2.85	0.00	0.00	0.00	0.00	271	String#== (ruby_runtime:0}
2.65	0.01	0.00	0.00	0.01	14	IRB::SLex::Node#match_io(d1) (/opt/ruby-enterprise/lib/ruby/1.8/irb/slex.rb:204}
2.55	0.02	0.00	0.00	0.02	20	RubyLex#token (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:279}
2.26	0.00	0.00	0.00	0.00	100	Kernel#puts (ruby_runtime:0}
2.23	0.00	0.00	0.00	0.00	20	Array#& (ruby_runtime:0}
1.72	0.02	0.00	0.00	0.02	2	RubyLex#lex (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:262}
1.71	0.02	0.00	0.00	0.02	20	IRB::SLex#match (/opt/ruby-enterprise/lib/ruby/1.8/irb/slex.rb:70}
1.65	0.00	0.00	0.00	0.00	1	Integer#times (ruby_runtime:0}
1.60	0.00	0.00	0.00	0.00	163	Fixnum#+ (ruby_runtime:0}
1.56	0.00	0.00	0.00	0.00	167	Kernel#hash (ruby_runtime:0}
1.53	0.00	0.00	0.00	0.00	147	Array#empty? (ruby_runtime:0}
1.24	0.00	0.00	0.00	0.00	20	RubyLex#peek (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:180}
1.24	0.00	0.00	0.00	0.00	44	<Class::RubyLex>#debug? (/opt/ruby-enterprise/lib/ruby/1.8/irb/ruby-lex.rb:34}

Rails command line

USAGE

```
script/performance/profiler 'Person.expensive_method(10)' [times] [flat|graph|graph_html]
```

EXAMPLE

```
script/performance/profiler 'Item.all'
```

Performance Testing: Measurement

Benchmark standard library

```
require 'benchmark'

puts Benchmark.measure { "a"*1_000_000 }
# 1.166667    0.050000    1.216667 ( 0.571355)

n = 50000
Benchmark.bm do |x|
  x.report { for i in 1..n; a = "1"; end }
  x.report { n.times do ; a = "1"; end }
  x.report { 1.upto(n) do ; a = "1"; end }
end
```

```
#          user      system      total      real
#  1.033333  0.016667  1.016667 ( 0.492106)
#  1.483333  0.000000  1.483333 ( 0.694605)
#  1.516667  0.000000  1.516667 ( 0.711077)
```

Rails command line

USAGE

script/performance/benchmark [times] 'Person.expensive_way' 'Person.another_expensive_way' ...

EXAMPLE

script/performance/benchmark 10 'Item.all' 'CouchItem.all'

Rails helper methods

Creating report in your log file

Model

```
Project.benchmark("Creating project") do
  project = Project.create("name" => "stuff")
  project.create_manager("name" => "David")
  project.milestones << Milestone.find(:all)
end
```

Controller

```
def process_projects
  self.class.benchmark("Processing projects") do
    Project.process(params[:project_ids])
    Project.update_cached_projects
  end
end
```

View

```
<% benchmark("Showing projects partial") do %>
  <%= render :partial => @projects %>
<% end %>
```

Performance Test cases

a special type of integration tests

- `script/generate performance_test` welcome
- `rake test:benchmark` (will run 4 times)
- `rake test:profile` (will run 1 time)

Performance Test cases

Example

```
require 'test_helper'
require 'performance_test_help'

class WelcomeTest < ActionController::PerformanceTest
  # Replace this with your real tests.
  def test_homepage
    get '/'
  end
end
```

rake test:benchmark

Started

BrowsingTest#test_homepage (7 ms warmup)

 wall_time: 2 ms

 memory: 414.53 KB

 objects: 2256

 gc_runs: 0

 gc_time: 0 ms

.WelcomeTest#test_homepage (4 ms warmup)

 wall_time: 2 ms

 memory: 414.53 KB

 objects: 2256

 gc_runs: 0

 gc_time: 0 ms

.

Finished in 0.908874 seconds.

10 tests, 0 assertions, 0 failures, 0 errors

rake test:profile

Started

BrowsingTest#test_homepage (6 ms warmup)

process_time: 12 ms

memory: 771.14 KB

objects: 2653

.WelcomeTest#test_homepage (4 ms warmup)

process_time: 12 ms

memory: 771.14 KB

objects: 2653

.

Finished in 3.49321 seconds.

6 tests, 0 assertions, 0 failures, 0 errors

Generic Tools

(black-box)

- `httperf`
- `ab` - Apache HTTP server benchmarking tool

How fast can this server serve requests?

- Use web server to serve static files as baseline measurement
- Do not run from the same server (I/O and CPU)
- Run from a machine as close as possible

You need know basic statistics

- compare not just their means but their standard deviations and confidence intervals as well.
- Approximately 68% of the data points lie within one standard deviation of the mean
- 95% of the data is within 2 standard deviation of the mean

httperf example

```
httperf --server localhost --port 3000 --uri / --num-conns 10000
```

```
httperf --client=0/1 --server=localhost --port=3000 --uri=/ --send-buffer=4096 --recv-buffer=16384 --num-conns=10000 --num-calls=1
```

```
httperf: warning: open file limit > FD_SETSIZE; limiting max. # of open files to FD_SETSIZE  
Maximum connect burst length: 1
```

```
Total: connections 10000 requests 10000 replies 10000 test-duration 18.373 s
```

```
Connection rate: 544.3 conn/s (1.8 ms/conn, <=1 concurrent connections)
```

```
Connection time [ms]: min 0.1 avg 1.8 max 4981.7 median 0.5 stddev 50.8
```

```
Connection time [ms]: connect 0.7
```

```
Connection length [replies/conn]: 1.000
```

```
Request rate: 544.3 req/s (1.8 ms/req)
```

```
Request size [B]: 87.0
```

```
Reply rate [replies/s]: min 55.0 avg 558.3 max 830.7 stddev 436.4 (3 samples)
```

```
Reply time [ms]: response 0.7 transfer 0.4
```

```
Reply size [B]: header 167.0 content 3114.0 footer 0.0 (total 3281.0)
```

```
Reply status: 1xx=0 2xx=10000 3xx=0 4xx=0 5xx=0
```


```
CPU time [s]: user 3.24 system 14.09 (user 17.7% system 76.7% total 94.4%)
```

```
Net I/O: 1790.1 KB/s (14.7*10^6 bps)
```

```
Errors: total 0 client-timo 0 socket-timo 0 connrefused 0 connreset 0
```

1 sample need 5s
30 samples will be accurate

Writing Efficient Ruby Code



Writing Efficient Ruby Code Tips

- Instance variable faster than accessor
- Interpolated string faster than + operator
- In-Place updates
- Module and class definition scope only execute once
- Caching Data in Instance or Class Variables
- Useless .nil?
- Unnecessary block parameter &block
- More...
 - <http://www.igvita.com/2008/07/08/6-optimization-tips-for-ruby-mri/>
 - <http://ihower.tw/blog/archives/1691>
 - <http://en.oreilly.com/rails2009/public/schedule/detail/8680>

Interpolated Strings

```
s = "#{a} #{b} #{c}"
```

is faster than

```
s = a.to_s + b.to_s + c.to_s
```


In-Place Updates

Class	Copying	Destructive

String	<code>#+</code>	<code>#<<</code>
String	<code>#sub</code>	<code>#sub!</code>
String	<code>#gsub</code>	<code>#gsub!</code>
Hash	<code>#merge</code>	<code>#merge!</code>
Array	<code>#+</code>	<code>#concat</code>
Array	<code>#map</code>	<code>#map!</code>
Array	<code>#compact</code>	<code>#compact!</code>
Array	<code>#uniq</code>	<code>#uniq!</code>
Array	<code>#flatten</code>	<code>#flatten!</code>

Coding Variable Caching (I)

```
def capital_letters  
  @capital_lettters ||= ("A".."Z").to_a  
end
```

or

```
@@capital_lettters ||= ("A".."Z").to_a  
def capital_letters  
  @@capital_lettters  
end
```

Coding Variable Caching (2)

```
def actions
  unless @actions
    # do something expensive
    @actions = expr
  end
```

```
  @actions
end
```

or

```
def actions
  @actions ||=
    begin
      # do something expensive
      expr
    end
end
```

ActiveSupport::Memoizable

```
def total_amount  
  # expensive calculate  
end
```

```
def total_income  
  # expensive calculate  
end
```

```
def total_expense  
  # expensive calculate  
end
```

```
extend ActiveSupport::Memoizable  
memoize :total_amount, :total_income, :total_expense
```

Method cache

- Ruby use cached method before method lookup
- Avoid these methods at runtime, it will clear cache.
 - `def / undef`
 - `Module#define_method`
 - `Module#remove_method`
 - `alias / Module#alias_method`
 - `Object#extend`
 - `Module#include`
 - `public/private/protected/module_function`

```
require 'benchmark'

class C
  def m; end
end

module H
end

puts Benchmark.measure {
  i = 0
  while i < 100000
    i+=1
    l = C.new
    # l.extend H
    l.m; l.m; l.m; l.m;
    l.m; l.m; l.m; l.m;
  end
}
```

```
require 'benchmark'
```

```
class C  
  def m; end  
end
```

```
module H  
end
```

```
puts Benchmark.measure {  
  i = 0  
  while i < 100000  
    i+=1  
    l = C.new  
    l.extend H  
    l.m; l.m; l.m; l.m;  
    l.m; l.m; l.m; l.m;  
  end  
}
```

```
require 'benchmark'

class C
  def m; end
end

module H
end

puts Benchmark.measure {
  i = 0
  x = C.new
  while i < 100000
    i+=1
    l = C.new
    x.extend H # Extend on an unrelated object!
    l.m; l.m; l.m; l.m;
    l.m; l.m; l.m; l.m;
  end
}
```



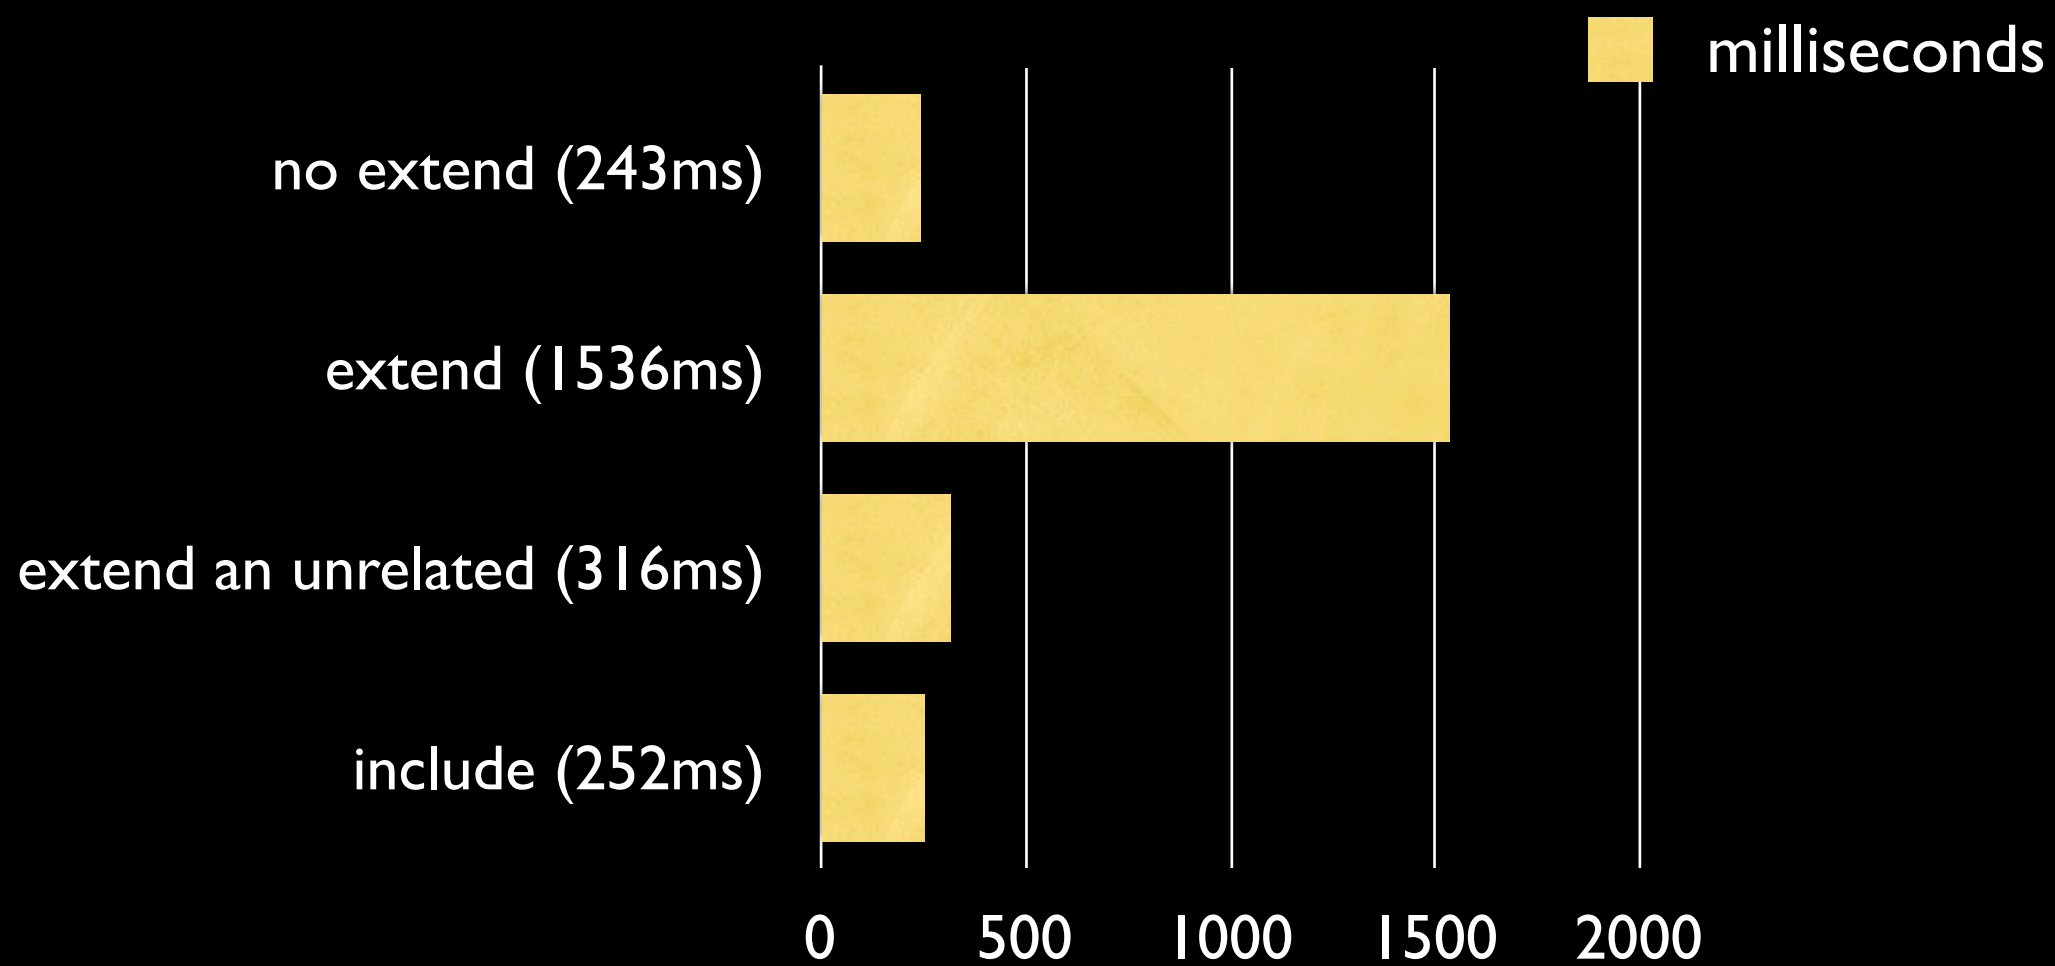
```
require 'benchmark'
```

```
class C  
  def m; end  
end
```

```
module H  
end
```

```
class MyC < C  
  include H  
end
```

```
puts Benchmark.measure {  
  i = 0  
  while i < 100000  
    i+=1  
    l = MyC.new  
    l.m; l.m; l.m; l.m;  
    l.m; l.m; l.m; l.m;  
  end  
}
```

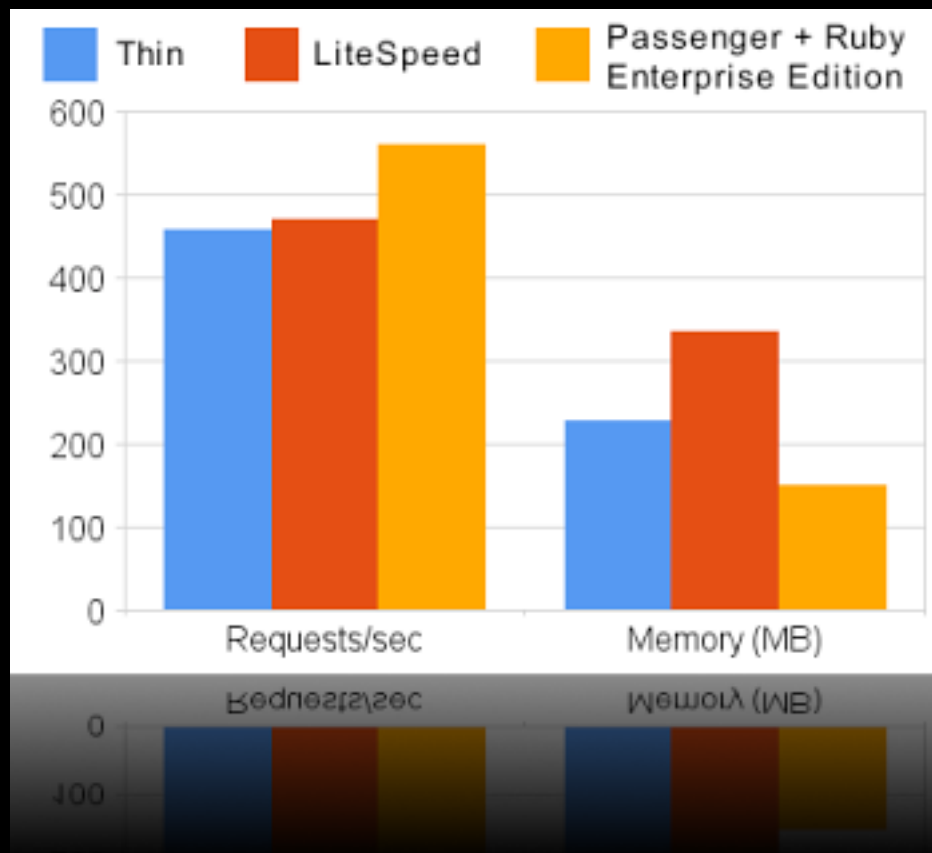


(MacBook Pro 2.2G)

Constant Caching

- Don't redefined constants at runtime
- Don't define new constants frequently

Use Ruby Enterprise Edition (REE) or Ruby 1.9



Ruby is slow?

- Language Micro-benchmarks != performance in complex systems
- Other factors:
 - application architecture
 - the ability to leverage higher-level
(Simplify things which may be complex to implement in other languages.
=> No code is faster than no code.)
 - Rails is faster than many PHP frameworks
<http://avnetlabs.com/php/php-framework-comparison-benchmarks>

Use faster Ruby Library

C Extension++

- XML parser

<http://nokogiri.org/>

- JSON parser

<http://github.com/brianmario/yajl-ruby/>

- CSV parser

<http://www.toastysapps.com/excelsior/>

- HTTP client

<http://github.com/pauldix/typhoeus>

- Date

<http://github.com/rtomayko/date-performance>

Caching

About Caching

- Ugly if you cache everywhere
- More bugs and tough to debug
including stale, data, inconsistent data, timing-based bugs.
- Complicated: expire, security
- Limit your user interface options

Cache Store

```
ActionController::Base.cache_store = :mem_cache_store, "localhost"
```

```
ActionController::Base.cache_store = :compressed_mem_cache_store, "localhost"
```

```
ActionController::Base.cache_store = :memory_store
```

```
ActionController::Base.cache_store = :synchronized_memory_store
```

```
ActionController::Base.cache_store = :file_store, "/path/to/cache/directory"
```

```
ActionController::Base.cache_store = :drb_store, "druby://localhost:9192"
```

View Caching

- Page Caching
- Action Caching
- Fragment Caching

Page Caching

```
class ProductsController < ActionController  
  caches_page :index  
  
  def index; end  
  
end
```

Action Caching

```
class ProductsController < ActionController

  before_filter :authenticate, :only => [ :edit, :create ]
  caches_action :edit

  def index; end

  def create
    expire_page :action => :index
    expire_action :action => :edit
  end

  def edit; end

end
```

Fragment Caching

```
<% cache(:key =>  
  ['all_available_products', @latest_product.created_at].join(':')) do %>  
  All available products:  
<% end %>
```

```
expire_fragment(:key =>  
  ['all_available_products', @latest_product.created_at].join(':'))
```

Sweepers

```
class StoreSweeper < ActionController::Caching::Sweeper
  # This sweeper is going to keep an eye on the Product model
  observe Product

  # If our sweeper detects that a Product was created call this
  def after_create(product)
    expire_cache_for(product)
  end

  # If our sweeper detects that a Product was updated call this
  def after_update(product)
    expire_cache_for(product)
  end

  # If our sweeper detects that a Product was deleted call this
  def after_destroy(product)
    expire_cache_for(product)
  end

  private
  def expire_cache_for(record)
    # Expire the list page now that we added a new product
    expire_page(:controller => '#{record}', :action => 'list')

    # Expire a fragment
    expire_fragment(:controller => '#{record}',
      :action => 'recent', :action_suffix => 'all_products')
  end
end
```

Caching yourself

```
Rails.cache.read("city")    # => nil  
Rails.cache.write("city", "Duckburgh")  
Rails.cache.read("city")    # => "Duckburgh"
```

```
Rails.cache.fetch("#{id}-data") do  
  Book.sum(:amount, :conditions => { :category_id => self.category_ids } )  
end
```

Use Memcached

- Free & open source, high-performance, distributed memory object caching system
- an in-memory key-value store for small chunks of arbitrary data (strings, objects) from results of database calls, API calls, or page rendering.

Memcached

- Key: 256 characters
- Data: 1mb
- SET/ADD/REPLACE/GET operators
- NOT persistent data store
- caching “noreply” principle

Caching secret

- Key naming
- Expiration

Caching Expire

- expire it after create/update/delete
 - race condition? lock it first.
- reset it after update
 - race condition first time? lock it first.
- set expire time
 - race condition? proactive cache refill

Using memcached



SQL and ActiveRecord


ORM is a high-level library that it's easy to forget about efficiency until it becomes a problem.

<http://www.engineyard.com/blog/2009/thats-not-a-memory-leak-its-bloat/>

Scaling Rails - On The Edge x Scaling Rails - On The Edge x Scaling Rails - On The Edge x git noahd1's oink at master - C x That's Not a Memory Leak, I x

← → ↺ ⌂ ☆ <http://www.engineyard.com/blog/2009/thats-not-a-memory-leak-its-bloat/> ▶ 📄 🔧

🌐 iGoogle 📄 Read Later 📄 Instapaper 🌐 Blog 🐦 Twitter 📘 Facebook 📄 Plurk 📄 Tumblr 📄 Yahoo! 📄 PlanetRoR 📄 git GitHub 📄 git Gist » 📁 Other Bookmarks



Sales: (866) 518-YARD [Login to Cloud](#)

Cloud Support Services Technology Blog Company



That's Not a Memory Leak, It's Bloat

By [Sudara Williams](#) | September 3rd, 2009 at 10:09AM

Our Rails customers often run into memory issues. The most frequent cause these days is what we in Support dub 'bloated mongrels.'

To be fair, bloat has absolutely nothing to do with mongrel itself, which is a solid and fine piece of work. You can run into this problem just as easily with thin, passenger, etc. Changing to a different server will not save you, as the root cause is not the server, but the code the server is running for you.

A real true-blooded memory leak is rare in comparison to the occurrence of bloating Rails instances. If your mongrels (or thins, or passenger instances) are suddenly sporting 100MB or more of extra weight, look no further: we've got the diet plan for you!

What Is Bloat?

In short: you are loading in too much. Too much what, you ask? Why it's too much ActiveRecord!

Bloat is *easily* identifiable. Last week, your mongrels were at 110MB, but after a new feature or two and a bit of 'optimization'.... well, lets just say that you'd have trouble fitting one on a CD. It's not always *that* dramatic (probably the average size of bloated mongrels are 200-300MB), but basically the mongrels are 2-5x larger than they should be, or spike in size suddenly after a certain subset of requests.

Detecting Bloat

The easiest way to detect bloat is to watch the Application Server process size. [New Relic](#), for example, will show you combined memory usage. You could watch it live with "top" on your



News from Engine Yard and the voices of our people.



The Engine Yard Newsletter

Get the latest Ruby and Rails tips, and news about Engine Yard Cloud.



Engine Yard Cloud Daily Demo

Monday - Friday

The next demo will begin in:

5 hrs, 1 min, 53 sec

Watch a 15-minute live demonstration of

N+1 Queries

```
# model
class User < ActiveRecord::Base
  has_one :car
end

class Car < ActiveRecord::Base
  belongs_to :user
end

# your controller
def index
  @users = User.paginate( :page => params[:page], :per_page => 20 )
end

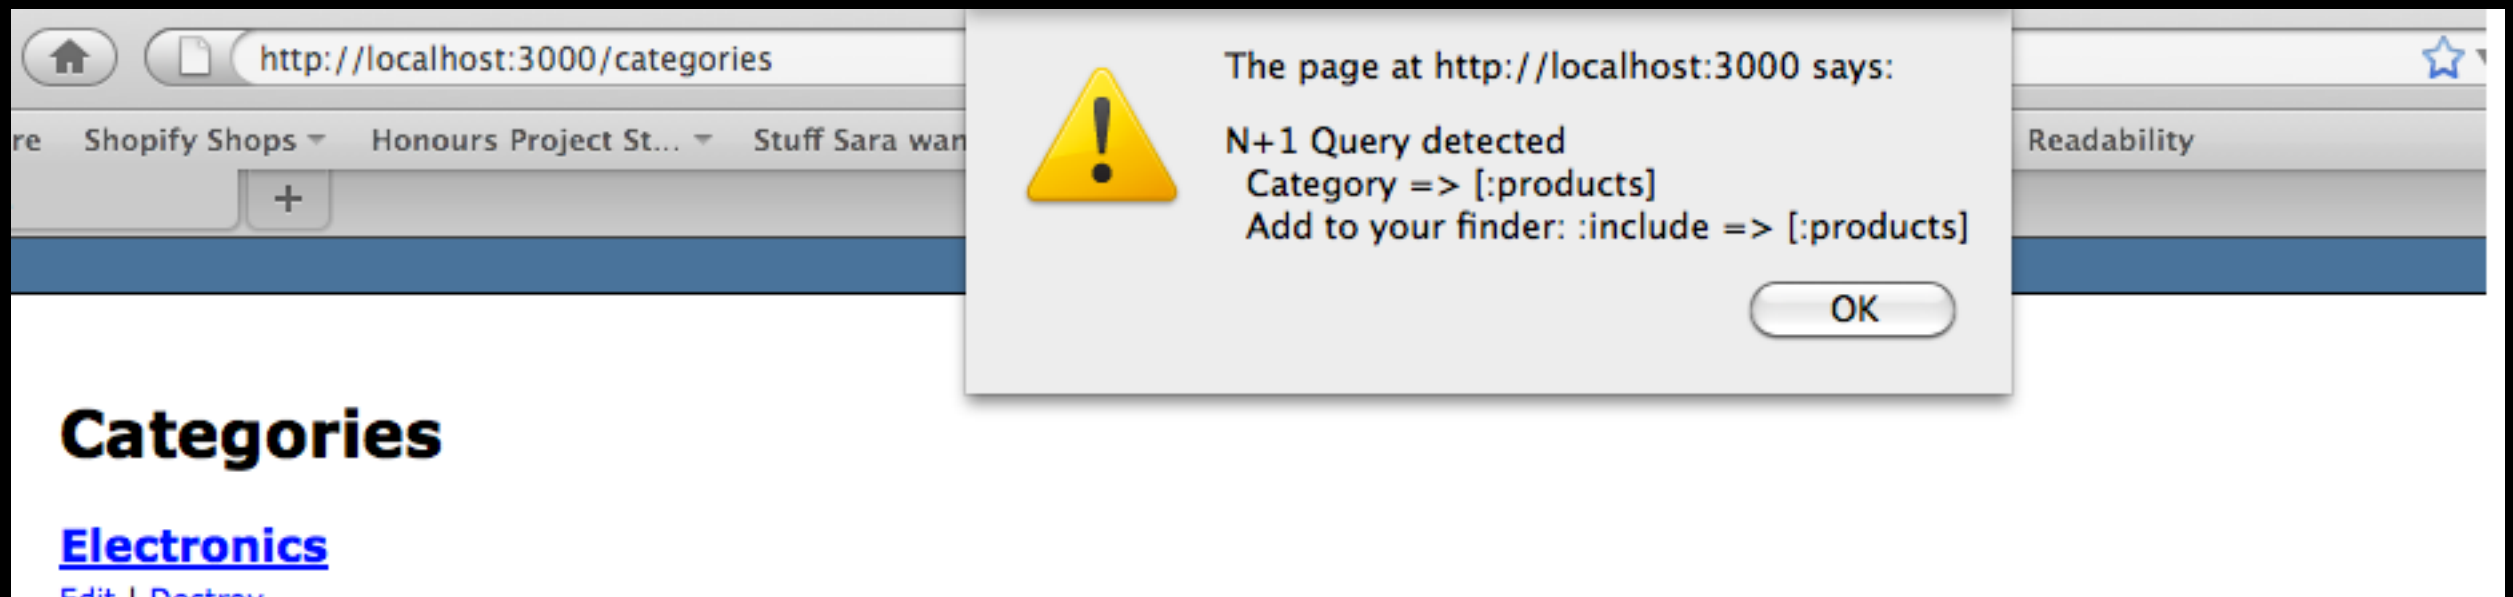
# view
<% @users.each do |user| %>
  <%= user.car.name %>
<% end %>
```

Add :include

```
# your controller
def index
  @users = User.paginate( :include => :car, :page => params
[:page], :per_page => 20 )
end
```


Bullet plugin

- <http://github.com/flyerhzm/bullet>
Help you reduce the number of queries with alerts (and growl).



Missing indexing

- Foreign key indexes
- Columns that need to be sorted
- Lookup fields
- Columns that are used in a GROUP BY
- http://github.com/eladmear/rails_indexes
Rake tasks to find missing indexes.

Only select you need

```
Event.find(:all, :select => "id, title, description")
```

```
class User < ActiveRecord::Base  
  named_scope :short, :select => "id, name, email"  
end
```

```
User.short.find(:all)
```

- <http://github.com/methodmissing/scrooge>
SQL query optimizer, so you query for only
what your page needs.

Replace :include to :join for some case

```
Group.find(:all, :include => [ :group_memberships ], :conditions =>  
[ "group_memberships.created_at > ?", Time.now - 30.days ] )
```

you can replace :include to :join

```
Group.find(:all, :joins => [ :group_memberships ], :conditions =>  
[ "group_memberships.created_at > ?", Time.now - 30.days ] )
```

Batch finding

```
Article.find_each do |a|  
  # iterate over all articles, in chunks of 1000 (the default)  
end
```

```
Article.find_each(:conditions => { :published => true }, :batch_size => 100 ) do |a|  
  # iterate over published articles in chunks of 100  
end
```

```
Article.find_in_batches do |articles|  
  articles.each do |a|  
    # articles is array of size 1000  
  end  
end
```

```
Article.find_in_batches(batch_size => 100 ) do |articles|  
  articles.each do |a|  
    # iterate over all articles in chunks of 100  
  end  
end
```

Transaction for group operations

```
my_collection.each do |q|  
  Quote.create({:phrase => q})  
end
```

```
# Add transaction  
Quote.transaction do  
  my_collection.each do |q|  
    Quote.create({:phrase => q})  
  end  
end
```

SQL query planner

- EXPLAIN keyword
- http://github.com/dsboulder/query_reviewer

SQL CRITICAL

Total queries: **3** Total time: **0.051s** MySQL Database Time: **0.002s**
With warnings: **2** (66%)
Without warnings: **1** (33%)
Type: **3** SELECTs

2 Errors:

0.001s Table products: Query Type: all, Key: (blank) [MSG](#) [SQL](#) [EXPLN](#) [TRACE](#) [IGNR](#)
Full table scan
No index was used here. In this case, that meant scanning 221 rows.


0.001s Table categories: Key: (blank) [MSG](#) [SQL](#) [EXPLN](#) [TRACE](#) [IGNR](#)

Safe queries:

0.001s SELECT SQL_NO_CACHE * FROM `categories` WHERE (`categories`.`id` = N) [EXPLN](#) [TRACE](#)

[Disable analysis report](#) on next page load and from now on.Sort by: [Name](#) | [Price](#)

SQL best practices



Use Full-text search engine

- Sphinx and thinking_sphinx plugin
- Ferret gem and acts_as_ferret

Use Constant for domain data

```
class Rating < ActiveRecord::Base
```

```
  G = Rating.find_by_name('G')
```

```
  PG = Rating.find_by_name('PG')
```

```
  R = Rating.find_by_name('R')
```

```
  #....
```

```
end
```

```
Rating::G
```

```
Rating::PG
```

```
Rating::R
```

Counter cache

```
class Topic < ActiveRecord::Base
  has_many :posts
end

class Posts < ActiveRecord::Base
  belongs_to :topic, :counter_cache => true
end

@topic.posts.size
```

Store Your Reports

- Aggregate reports via cron and rake

AR Caching Plugins

- <http://github.com/nkallen/cache-money>
- <http://github.com/fauna/interlock>
- http://github.com/defunkt/cache_fu
- Need careful to go to these solution, because it's very intrusive.

Consider NoSQL data store

- **Key-value stores for high performance**
Redis, Tokyo Cabinet, Flare
- **Document stores for huge storage**
MongoDB, CouchDB
- **Record store for high scalability and availability**
Cassandra, HBase, Voldemort

Use key-value store from now

- Redis, Tokyo Cabinet are very very fast
- Avoid touching the RDBMS when storing non-critical data, hit count, download count, online users count...etc

Moneta

- a unified interface to key/value stores
- <http://github.com/wycats/moneta/tree>

moneta supports:

- File store for xattr
- Basic File Store
- Memcache store
- In-memory store
- The xattrs in a file system
- DataMapper
- S3
- Berkeley DB
- Redis
- SDBM
- Tokyo
- CouchDB

moneta API:

- #[](key)
- #[]=(key, value)
- #delete(key)
- #key?(key)
- #store(key, value, options)
- #update_key(key, options):
- #clear

example

```
begin
  # for developer has tokyo cabinet
  STORE = Moneta::Tyrant.new( :host => 'localhost', :port => 1978 )
rescue
  # for developer has not tokyo cabinet
  STORE = Moneta::BasicFile.new( :path => "tmp" )
end

STORE["mykey"] = { :foo => 111 , :bar => 222 }
```

Rails Metal

a subset of Rack middleware

- 2~3x faster than a controller, because it bypasses routes and controller.
- APIs and anything which need not ActionView

Use web server or CDN for static file

- Web server is 10x faster than your application server
- Set `:x_sendfile` to true if you use Apache `mod_xsendfile` or Lighttpd

Web performance client-side analysis

- <http://developer.yahoo.com/yslow/>
- <http://code.google.com/speed/page-speed/>

Web performance


Rules 14

- Make Fewer HTTP Requests
- Use a Content Delivery Network
- Add an Expires Header
- Gzip Components
- Put Stylesheets at the Top
- Put Scripts at the Bottom
- Avoid CSS Expressions
- Make JavaScript and CSS External
- Reduce DNS Lookups
- Minify JavaScript
- Avoid Redirects
- Remove Duplicates Scripts
- Configure ETags
- Make Ajax Cacheable

14 Steps to Faster-Loading Web Sites

High Performance Web Sites

*Essential Knowledge
for Front-End Engineers*



O'REILLY®

*Steve Souders
Foreword by Nate Koechley*

Essential Knowledge for Frontend Engineers



Even Faster Web Sites

O'REILLY®

Steven Souders

Use external programs

```
def thumbnail(temp, target)
  system("/usr/local/bin/convert #{escape(temp)} -resize
48x48! #{escape(target)}")
end
```

Write inline C/C++ code

- **RubyInline:** Write foreign code within ruby code
<http://rubyinline.rubyforge.org/RubyInline/>
- **Rice:** Ruby Interface for C++ Extensions
<http://rice.rubyforge.org/>
- **Ruby-FFI:** a ruby extension for programmatically loading dynamic libraries
<http://github.com/ffi/ffi>

Reference

- Advanced Rails Chap.6 Performance (O'Reilly)
- Rails Rescue Handbook
- Writing Efficient Ruby Code (Addison-Wesley)
- Ruby on Rails Code Review (Peepcode)
- Rails 2 Chap. 13 Security and Performance Enhancements (friendsof)
- Deploying Rails Application Chap.9 Performance (Pragmatic)
- http://guides.rubyonrails.org/caching_with_rails.html
- http://guides.rubyonrails.org/performance_testing.html
- <http://railslab.newrelic.com/scaling-rails>
- <http://antoniochangiano.com/2007/02/10/top-10-ruby-on-rails-performance-tips/>
- <http://www.engineyard.com/blog/2009/thats-not-a-memory-leak-its-bloat/>
- <http://jstorimer.com/ruby/2009/12/13/essential-rails-plugins-for-your-inner-dba.html>
- <http://asciicasts.com/episodes/161-three-profiling-tools>
- <http://robots.thoughtbot.com/post/163627511/a-grand-piano-for-your-violin>

So, how to scale?

- Rails performance (AR...etc how to ruby code)
- Web performance (yslow related)
- Asynchrony Processing (Message queue)
- HTTP Reverse Proxy Caching
- Partition Component using SOA
- Distributed Filesystem/Database

TODO (maybe next time)

- NoSQL: Key-value data store
- Front-end web performance
- HTTP reverse proxy caching

The End

感謝聆聽